

Bangladesh University of Engineering and Technology

EEE 6608

Machine Learning & Pattern Recognition

Final assignment: Implementing K-means clustering algorithm from
scratch

Submitted by,
Name: Md. Al-Imran Abir
Student ID: 0421062549

Date of Submission: April 17, 2022

1 Code:

```
1  # Libraries
2
3  import numpy as np
4  import pandas as pd
5
6  from sklearn.datasets import load_digits
7  from sklearn.metrics import calinski_harabasz_score, silhouette_score
8  from sklearn.cluster import KMeans
9
10 import matplotlib.pyplot as plt
11
12 #----- Load data-----
13
14 digits = load_digits()
15 features = digits.data
16 labels = digits.target
17 print(f"Features -> type:{type(features)}; shape:{features.shape}")
18 print(f"Labels   -> type:{type(labels)}; shape:{labels.shape}")
19
20 #-----
21 #-----PCA-----
22 #-----
23
24 # sample average
25 mu = np.mean(features, axis=0)
26 print(f"Shape of mean:{mu.shape}")
27
28 # mean centered sample
29 X = features - mu
30 # print(np.mean(X, axis=0))
31
32 # the covariance matrix
33 A = np.cov(X, rowvar=False)
34 print(f"shape of covariance matrix:{A.shape}")
35 # eigenvalues and right eigenvectors
36 eig_value, eig_vect = np.linalg.eig(A)
37
38 print(f"eigenValues: {eig_value.shape}; type: {type(eig_value)}")
39 print(f"eigenValues: {eig_vect.shape}; type: {type(eig_vect)}")
40
41
42 # find the value of K for which 95% of energy is retained
43 k = 1
```

```

44 energy = 0
45 while energy <= 0.95:
46     eig_value_k = eig_value[:k]
47     energy = eig_value_k.sum()/eig_value.sum()
48     k += 1
49     # print(energy)
50 print(f"k:{k} -> energy:{energy}")
51
52 # choose the first k eigen vectors
53 U_k = eig_vect[:, :k]
54 print(f"shape of projection matrix: {U_k.shape}")
55
56 # Projection
57 X_proj = X.dot(U_k)
58 print(f"projected matrix shape: {X_proj.shape}")
59
60
61 #-----
62 #----- K-means clustering-----
63 #-----
64
65 def compute_cost(X, centroids, cluster):
66     m = X.shape[0]
67     sum = 0
68     for i in range(m):
69         sum += (X[i] - centroids[int(cluster[i])]) @ (X[i] - centroids[int(cluster[i])]).T
70     cost = sum/m
71     return cost
72
73 def my_kmeans(X, k, seed=41):
74     diff = 1
75     np.random.seed(seed)
76     cluster_labels = np.zeros(X.shape[0]) #cluster label for each sample
77     rand_idx = np.random.choice(len(X), k, replace=False)
78     #Randomly choosing Centroids
79     centroids = X[rand_idx, :] #Step 1
80
81     while diff:
82         for i, row in enumerate(X):
83             min_dist = float('inf')
84             for idx, centroid in enumerate(centroids):
85                 d = np.sqrt((row - centroid) @ (row-centroid).T)
86                 if min_dist > d:
87                     min_dist = d
88                     cluster_labels[i] = idx
89     new_centroids = pd.DataFrame(X).groupby(by=cluster_labels).mean().values

```

```

90         if np.count_nonzero(centroids-new_centroids) == 0:
91             diff = 0
92         else:
93             centroids = new_centroids
94             cost = compute_cost(X, centroids, cluster_labels)
95             cal_score = calinski_harabasz_score(X, cluster_labels)
96             return centroids, cluster_labels, cost, cal_score
97
98
99 #-----
100 #----- With PCA -----
101 #-----
102
103 def plot_cost(cost, k):
104     plt.plot(range(2, k), cost)
105     plt.show()
106
107 #### Variation wrt No of Clusters
108
109 K = 15
110 cost = np.zeros(K-2)
111 sil_score = np.zeros(K-2)
112 cal_score = np.zeros(K-2)
113 centroids_lst = []
114 labels = np.zeros((K-1, X.shape[0]))
115
116 for k in range(2, K):
117     centroids, labels[k-2], cost[k-2], cal_score[k-2] = my_kmeans(X_proj, k)
118     centroids_lst.append(centroids)
119     sil_score[k-2] = silhouette_score(X_proj, labels[k-2])
120     print(f"k={k} -> cost:{cost[k-2]:.3f};
121           calinski:{cal_score[k-2]:.3f}; sil:{sil_score[k-2]:.3f}")
122
123 # print(f"silhouette_score:{sil_score}\nCost: {cost}\ncalinski_harabasz_score: {cal_score}")
124 plot_cost(cost, K)
125
126 ### Plot cost vs clusters
127 plt.figure()
128 plt.plot(range(2, K), cost)
129 plt.savefig("pca_cost_vs_clusters.png")
130 plt.xlabel("No of clusters")
131 plt.ylabel("Cost")
132 plt.show()
133
134 ### Plot Calinski-Harabasz Index vs clusters
135 plt.figure()

```

```

136 plt.plot(range(2, K), cal_score)
137 plt.xlabel("No of clusters")
138 plt.ylabel("Calinski-Harabasz Index")
139 plt.savefig("pca_cal_score_vs_clusters.png")
140 plt.show()
141
142 ### Plot Silhouette Coefficient vs clusters
143 plt.figure()
144 plt.plot(range(2, K), sil_score)
145 plt.xlabel("No of clusters")
146 plt.ylabel("Silhouette Coefficient")
147 plt.savefig("pca_cost_vs_sil_score.png")
148 plt.show()
149
150
151 # Variation wrt No of Samples-----
152
153
154 ki = 10
155 per = np.linspace(0.1, 1, num=10)
156 z = per.shape[0]
157 cost_sam = np.zeros(z)
158 sil_score_sam = np.zeros(z)
159 cal_score_sam = np.zeros(z)
160 centroids_lst_sam = []
161 labels_lst_sam = []
162 for i, j in enumerate(per):
163     idx = np.random.choice(range(X_proj.shape[0]), int(j*X_proj.shape[0]), replace=False)
164     X_per = X_proj[idx,:]
165     # print(i, j, X_per.shape[0])
166     centroids_sam, labels_sam, cost_sam[i], cal_score_sam[i] = my_kmeans(X_per, ki)
167     centroids_lst_sam.append(centroids_sam)
168     labels_lst_sam.append(labels_sam)
169     sil_score_sam[i] = silhouette_score(X_per, labels_lst_sam[i])
170     print(f"j={j:.2f} -> cost:{cost_sam[i]:.3f};
171           calinski:{cal_score_sam[i]:.3f}; sil:{sil_score_sam[i]:.3f}")
172
173 plt.plot(per, sil_score_sam)
174
175 ### Plot cost vs sample size
176 plt.figure()
177 plt.plot(per, cost_sam)
178 plt.savefig("pca_cost_vs_per_sample.png")
179 plt.xlabel("% of total sample")
180 plt.ylabel("Cost")

```

```

181 plt.show()
182
183 ### Plot Calinski-Harabasz Index vs sample size
184 plt.figure()
185 plt.plot(per, cal_score_sam)
186 plt.savefig("pca_cal_vs_per_sample.png")
187 plt.xlabel("% of total sample")
188 plt.ylabel("Calinski-Harabasz Index")
189 plt.show()
190
191 ### Plot Silhouette Coefficient vs sample size
192 plt.figure()
193 plt.plot(per, sil_score_sam)
194 plt.savefig("pca_sim_vs_per_sample.png")
195 plt.xlabel("% of total sample")
196 plt.ylabel("Silhouette Coefficient")
197 plt.show()
198
199 # Different Initialization-----
200
201 seeds = [21, 75, 84, 12, 51]
202 ls = len(seeds)
203 cost_in = np.zeros(ls)
204 cal_in = np.zeros(ls)
205 sil_in = np.zeros(ls)
206 for ij, seed in enumerate(seeds):
207     labels_in, cost_in[ij], cal_in[ij] = my_kmeans(X_proj, 10, seed=seed)[1:4]
208     sil_in[ij] = silhouette_score(X_proj, labels_in)
209     print(f" {seed} -> cost:{cost_in[ij]:.3f};
210           cal:{cal_in[ij]:.3f}; sil:{sil_in[ij]:.3f}")
211
212 #-----
213 #----- Without PCA -----
214 #-----
215
216 # Variation wrt No of Clusters -----
217
218 K_w = 15
219 cost_w = np.zeros(K_w-2)
220 sil_score_w = np.zeros(K_w-2)
221 cal_score_w = np.zeros(K_w-2)
222 centroids_lst_w = []
223 labels_w = np.zeros((K_w-1, X.shape[0]))
224
225 for k in range(2, K_w):

```

```

226     centroids_w, labels_w[k-2], cost_w[k-2], cal_score_w[k-2] = my_kmeans(X, k)
227     centroids_lst_w.append(centroids_w)
228     sil_score_w[k-2] = silhouette_score(X, labels_w[k-2])
229     print(f"k={k} -> cost:{cost_w[k-2]}; calinski:{cal_score_w[k-2]};
230           sil:{sil_score_w[k-2]}")
231
232     # print(f"silhouette_score:{sil_score}\nCost: {cost}\ncalinski_harabasz_score: {cal_score}")
233     plot_cost(cost_w, K)
234
235     plt.figure()
236     plt.plot(range(2, K_w), cost_w)
237     plt.savefig("no_pca_cost_vs_clusters.png")
238     plt.xlabel("No of clusters")
239     plt.ylabel("Cost")
240     plt.show()
241
242     plt.figure()
243     plt.plot(range(2, K_w), cal_score_w)
244     plt.savefig("no_pca_cal_vs_clusters.png")
245     plt.xlabel("% of total sample")
246     plt.ylabel("Calinski-Harabasz Index")
247     plt.show()
248
249     plt.figure()
250     plt.plot(range(2, K_w), sil_score_w)
251     plt.savefig("no_pca_sil_vs_clusters.png")
252     plt.xlabel("% of total sample")
253     plt.ylabel("Silhouette Coefficient")
254     plt.show()
255
256     # Variation wrt No of Samples-----
257
258     ki_w = 10
259     per_w = np.linspace(0.1, 1, num=10)
260     z_w = per_w.shape[0]
261     cost_sam_w = np.zeros(z_w)
262     sil_score_sam_w = np.zeros(z_w)
263     cal_score_sam_w = np.zeros(z_w)
264     centroids_lst_sam_w = []
265     labels_lst_sam_w = []
266     for i, j in enumerate(per_w):
267         idx_w = np.random.choice(range(X.shape[0]), int(j*X.shape[0]), replace=False)
268         X_per_w = X[idx_w,:]
269         # print(i, j, X_per_w.shape[0])
270         centroids_sam_w, labels_sam_w, cost_sam_w[i], cal_score_sam_w[i] = my_kmeans(X_per_w, ki_w)

```

```

271     centroids_lst_sam_w.append(centroids_sam_w)
272     labels_lst_sam_w.append(labels_sam_w)
273     sil_score_sam_w[i] = silhouette_score(X_per_w, labels_lst_sam_w[i])
274     print(f"j={j:.2f} -> cost:{cost_sam_w[i]:.3f};
275           calinski:{cal_score_sam_w[i]:.3f}; sil:{sil_score_sam_w[i]:.3f}")
276
277 plt.figure()
278 plt.plot(per_w, cost_sam_w)
279 plt.savefig("no_pca_cost_vs_per_sample.png")
280 plt.xlabel("% of total sample")
281 plt.ylabel("Cost")
282 plt.show()
283
284 plt.figure()
285 plt.plot(per_w, cal_score_sam_w)
286 plt.savefig("no_pca_cal_vs_per_sample.png")
287 plt.xlabel("% of total sample")
288 plt.ylabel("Calinski-Harabasz Index")
289 plt.show()
290
291 plt.figure()
292 plt.plot(per_w, sil_score_sam_w)
293 plt.savefig("no_pca_sil_vs_per_sample.png")
294 plt.xlabel("% of total sample")
295 plt.ylabel("Silhouette Coefficient")
296 plt.show()
297
298 # Different Initialization-----
299
300 seeds_w = [48, 14, 97, 62, 53]
301 ls_w = len(seeds_w)
302 cost_in_w = np.zeros(ls_w)
303 cal_in_w = np.zeros(ls_w)
304 sil_in_w = np.zeros(ls_w)
305 for ij, seed in enumerate(seeds_w):
306     labels_in_w, cost_in_w[ij], cal_in_w[ij] = my_kmeans(X, 10, seed=seed)[1:4]
307     sil_in_w[ij] = silhouette_score(X, labels_in_w)
308     print(f" {seed} -> cost:{cost_in_w[ij]:.3f};
309           cal:{cal_in_w[ij]:.3f}; sil:{sil_in_w[ij]:.3f}")

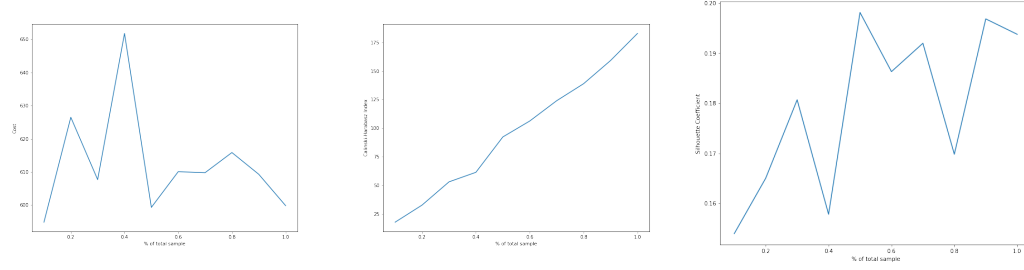
```


2 Results

2.1 With PCA

2.1.1 Variation of clustering performance with number of samples

Cost, Calinski-Harabasz Index, and Silhouette Coefficients while varying the sample size are shown in fig. 1a, fig. 1b, and fig. 1c respectively.

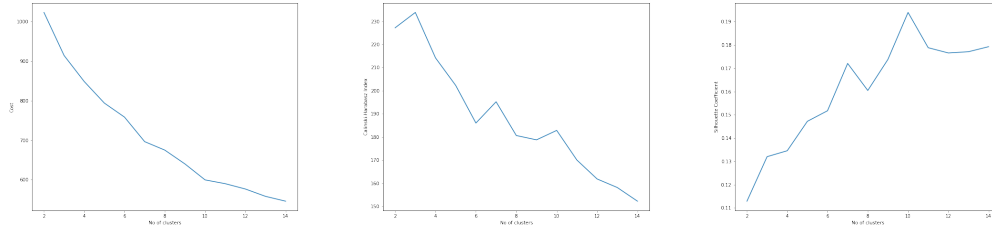


(a) Cost vs no of sample plot (b) Calinski-Harabasz Index vs (c) Silhouette Coefficient vs no of samples plot

Figure 1: Variation of clustering performance with number of samples (with PCA)

2.1.2 Variation of clustering performance with number of clusters K

Cost, Calinski-Harabasz Index, and Silhouette Coefficients while varying the no of clusters are shown in fig. 2a, fig. 2b, and fig. 2c respectively.



(a) Cost vs no of clusters plot (b) Calinski-Harabasz Index vs (c) Silhouette Coefficient vs no of clusters

Figure 2: Variation of clustering performance with number of clusters K (with PCA)

2.1.3 Variation of clustering performance with different initializations of the cluster centers

Different initialization was performed by setting the seed of NumPy pseudo-random number generator (`numpy.random.seed(seed)`) to different numbers. The corresponding performance metrics

are shown in table 1.

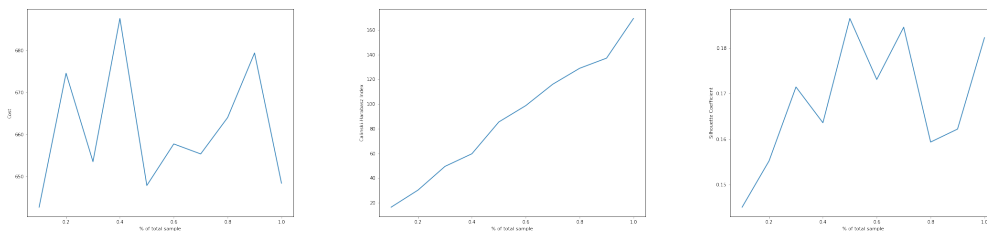
Table 1: Variation of clustering performance with different initializations of the cluster centers (with PCA)

Seed	Cost	Calinski-Harabasz Index	Silhouette Coefficient
21	603.331	180.672	0.195
75	619.554	170.742	0.182
84	643.371	157.071	0.160
12	624.812	167.634	0.179
51	602.507	181.190	0.198

2.2 Without PCA

2.2.1 Variation of clustering performance with number of samples

Cost, Calinski-Harabasz Index, and Silhouette Coefficients while varying the sample size are shown in fig. 3a, fig. 3b, and fig. 3c respectively.



(a) Cost vs no of samples plot (b) Calinski-Harabasz Index vs no of samples (c) Silhouette Coefficient vs no of samples

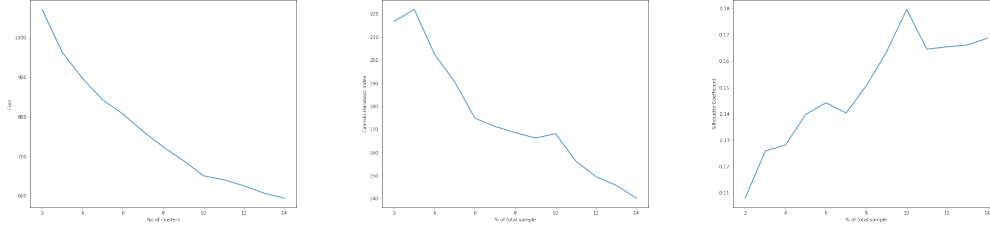
Figure 3: Variation of clustering performance with number of samples (without PCA)

2.2.2 Variation of clustering performance with number of clusters K

Cost, Calinski-Harabasz Index, and Silhouette Coefficients while varying the no of clusters are shown in fig. 4a, fig. 4b, and fig. 4c respectively.

2.2.3 Variation of clustering performance with different initializations of the cluster centers

Different initialization was performed by setting the seed of NumPy pseudo-random number generator (`numpy.random.seed(seed)`) to different numbers. The corresponding performance metrics are shown in table 2.



(a) Cost vs no of clusters plot (b) Calinski-Harabasz Index vs no of clusters (c) Silhouette Coefficient vs no of clusters

Figure 4: Variation of clustering performance with number of clusters K (without PCA)

Table 2: Variation of clustering performance with different initializations of the cluster centers (without PCA)

Seed	Cost	Calinski-Harabasz Index	Silhouette Coefficient
48	650.183	168.357	0.188
14	650.176	168.361	0.188
97	682.802	150.829	0.163
62	650.852	167.980	0.187
53	659.397	163.230	0.173

3 Discussion

In this project, the K-means clustering algorithm was implemented from scratch. Also, principal component analysis (PCA) was used for dimensionality reduction. We investigated the variation in the performances of the K-means algorithm with respect to different aspects.

3.1 Principal Component Analysis

We implemented PCA using the `numpy.linalg.eig()` function which computes the eigenvalues and right eigenvectors of a square matrix. As per the instruction, we also retained 95% energy in the reduced "k" dimensional space. The energy can be calculated by summing the corresponding eigenvalues.

The original dataset had 64 features. To capture 95% of the original energy, we needed only 30 principal features. These 30 features retained 95.48% of the total energy.

3.2 K-Means Clustering

Then we performed K-means clustering on both the given data samples and the reduced featured samples. We varied the no of clusters, sample size and observed three metrics, namely the loss function, Calinski-Harabasz Index and Silhouette Coefficient. We also observed the effect of initialization.

3.2.1 With PCA

At first, we performed K-means clustering on the reduced dimensional space. As we varied the size of the dataset (i.e. no of samples), we found that Calinski-Harabasz Index increases monotonically with no of samples (fig. 1b) and Silhouette Coefficient has an increasing trend (fig. 1c). But the cost function has no regular shape with respect to the no of sample (fig. 1a).

With the increase in no of clusters, we found that the cost function decreases monotonically as expected (fig. 2a). But there is no clear "knee" in the plot from which we may reach a conclusion about the optimal no of clusters. However, the Silhouette Coefficient (fig. 2c) is highest for 10 clusters and also the Calinski-Harabasz Index (fig. 2b) has a higher value at 10 than the neighbors. So, we may conclude that there are 10 clusters in the dataset. This is also supported by our prior knowledge about the dataset (the dataset has 10 different classes).

We also observe the effect of different initialization on the performance metrics (table 1). Depending on the initialization the performance metrics could vary in a large scale. However for our case, the performance metrics didn't vary that much.

3.2.2 Without PCA

Using the full feature space, we found also more or less the same results as with PCA (reduced dimensional space, fig. 3, fig. 4). The performance metrics didn't vary that much. This proves the fact that using PCA, we can achieve similar performances as using the main dataset with the added benefit of reduced computational resources.